



# Edges and binary images

Thursday, Sept 15, 2020

Richang Hong, Hefei University of Technology

COPYRIGHT@Kristen Grauman, UT Austin



# Outline

- Last time
  - Image gradients
  - Seam carving – gradients as “energy”
  - Gradients  $\rightarrow$  edges and contours
  - Template matching
    - Image patch as a filter
- Today
  - Binary image analysis
    - Blobs and regions



# Gradients -> edges



Primary edge detection steps:

1. Smoothing: suppress noise
2. Edge enhancement: filter for contrast
3. Edge localization

Determine which local maxima from filter output are actually edges vs. noise

- Threshold, Thin

# Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
  - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny');`
- `>>help edge`

# The Canny edge detector

---



original image (Lena)

# The Canny edge detector



---




norm of the gradient

# The Canny edge detector

---



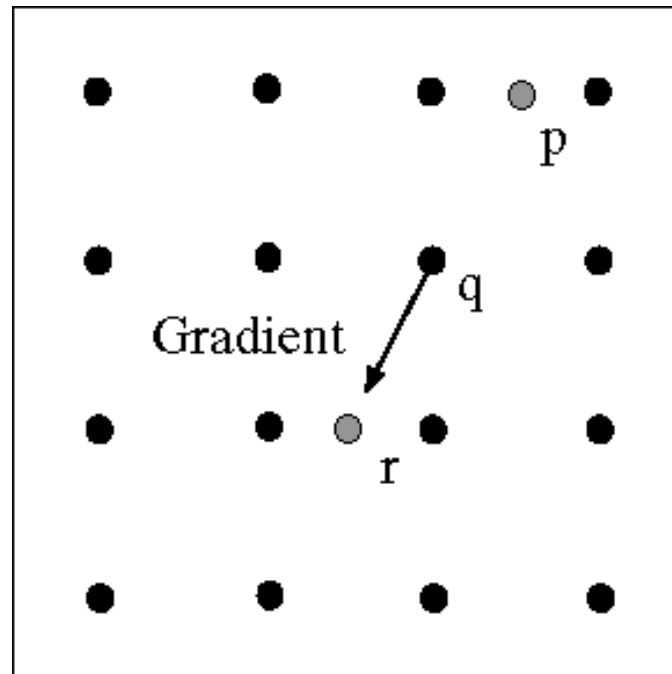
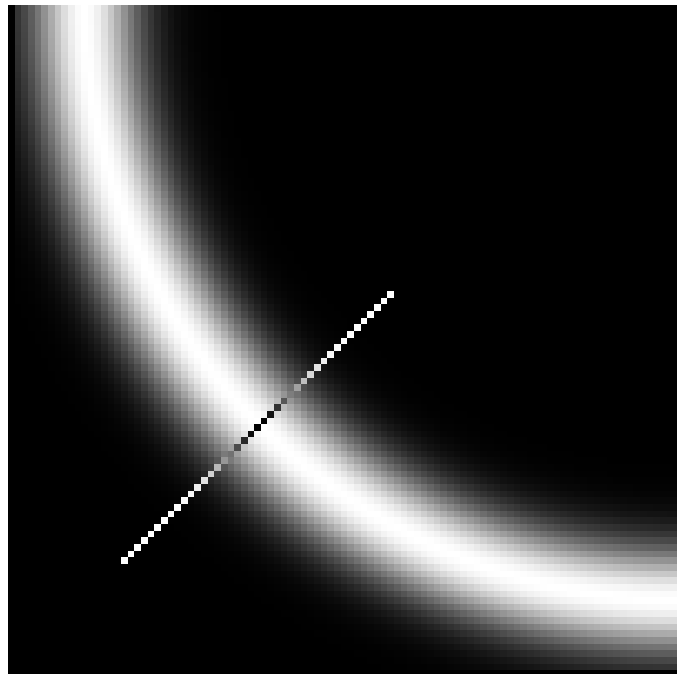
**Non-maximum suppression** :  
Thin wide "ridges"  
down to single  
pixel width



How to turn  
these thick  
regions of the  
gradient into  
curves?

# Non-maximum suppression

---



Check if pixel is local maximum along gradient direction,  
select single max across width of the edge

- requires checking interpolated pixels p and r



# The Canny edge detector

---



Problem:  
pixels along  
this edge  
didn't  
survive the  
thresholding

thinning  
(non-maximum suppression)

# Hysteresis thresholding


- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



Credit: James Hays

# Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them.



# Final Canny Edges



Credit: James Hays

# Recap: Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
  - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny');`
- `>>help edge`




# Low-level edges vs. perceived contours

---



image



human segmentation



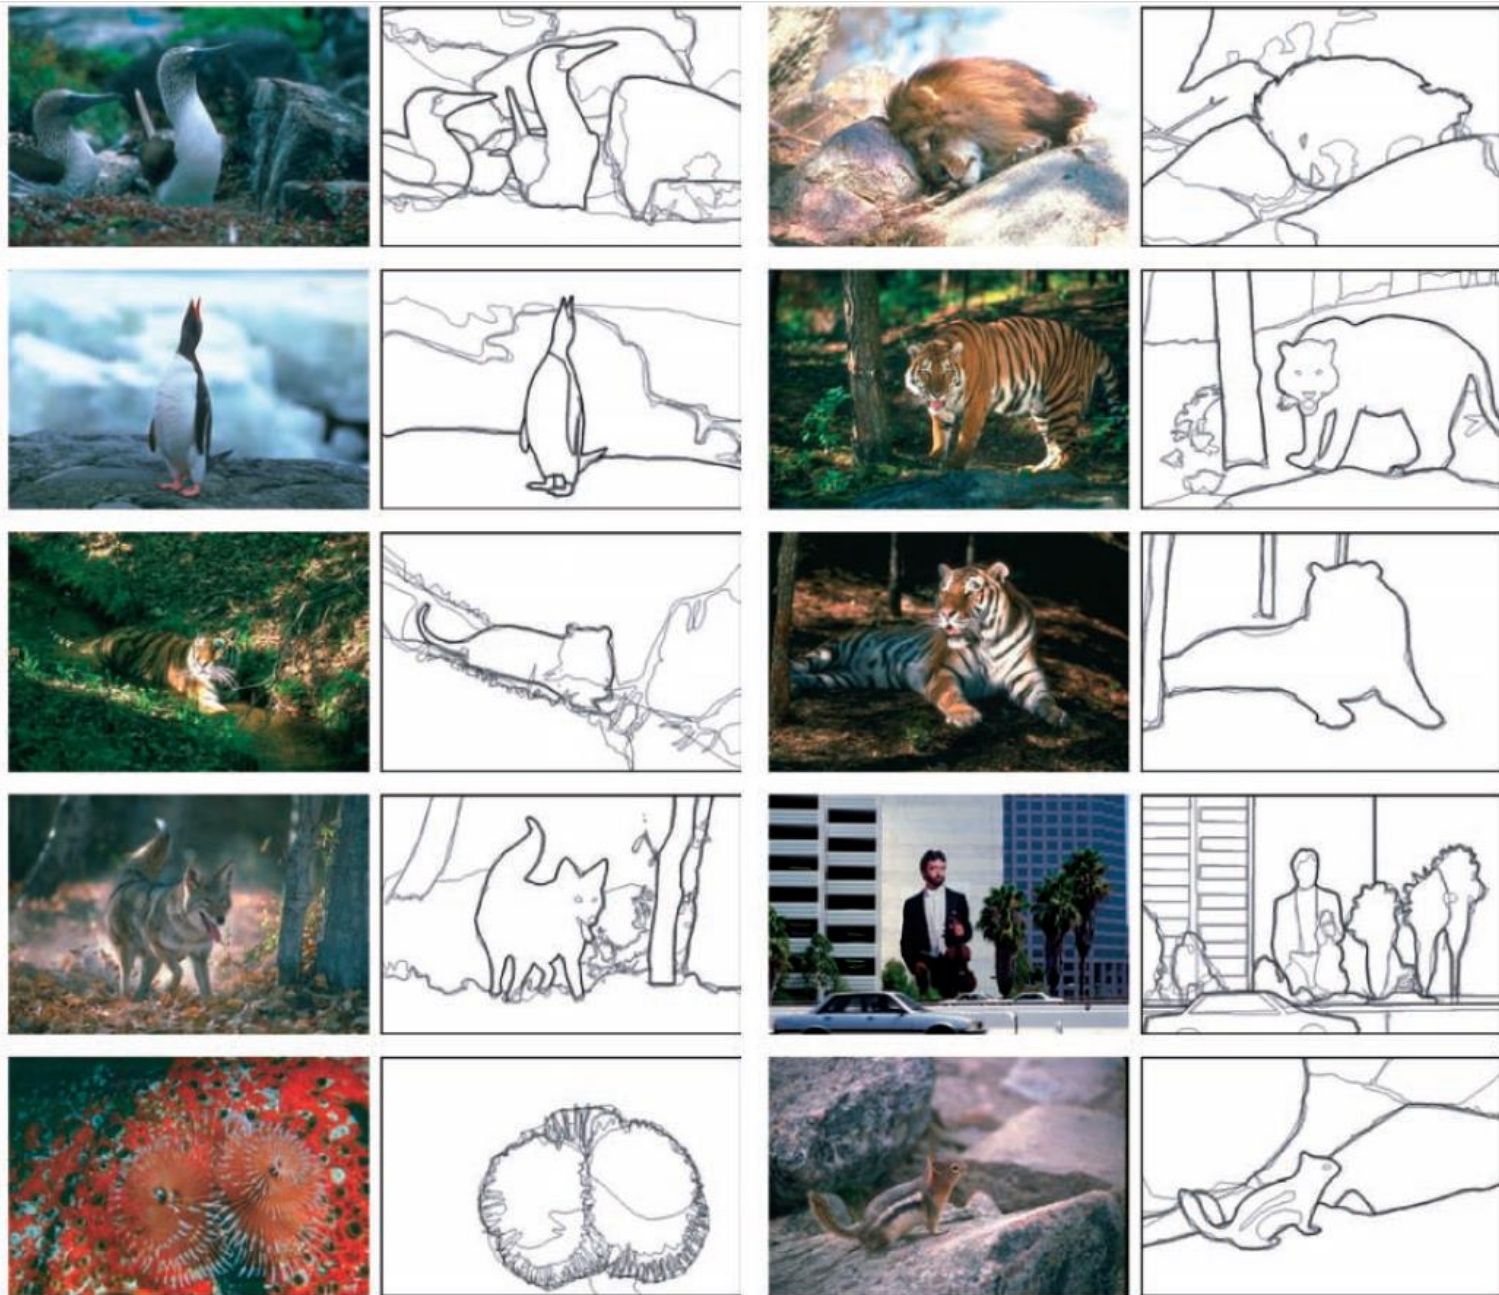
gradient magnitude



Berkeley segmentation database:

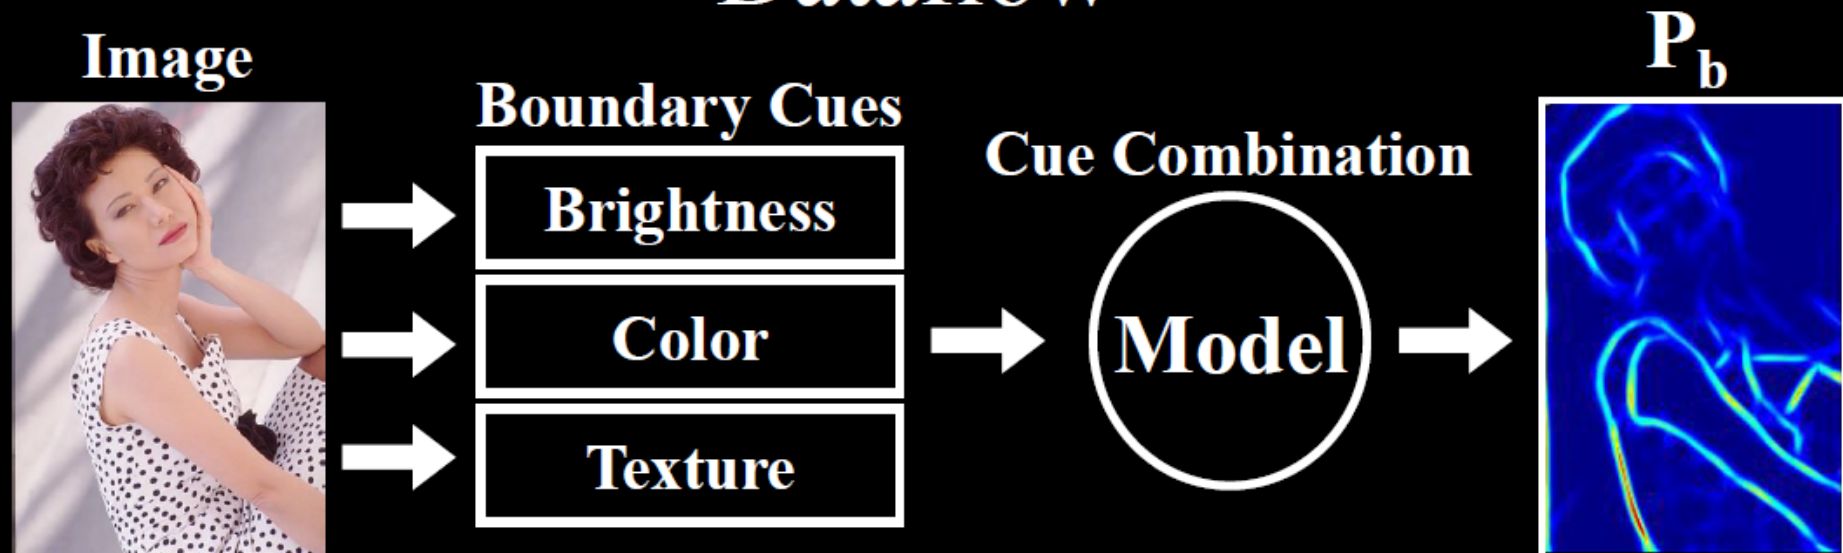
<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

Learn from humans which combination of features is most indicative of a “good” contour?





# Dataflow




Challenges: texture cue, cue combination


Goal: learn the posterior probability of a boundary  $P_b(x,y,\theta)$  from local information only




Image



BG+CG+TG



Human



# Recall: image filtering

- Compute a function of the local neighborhood at each pixel in the image
  - Function specified by a “filter” or mask saying how to combine values from neighbors.
- Uses of filtering:
  - Enhance an image (denoise, resize, etc)
  - Extract information (texture, edges, etc)
  - Detect patterns (template matching)



# Where's Waldo?



**Scene**



**Template**



# Where's Waldo?



**Template**


**Detected template**



# Where's Waldo?



**Detected template**



**Correlation map**

# Recap: Mask properties

- Smoothing

- Values positive
- Sum to 1  $\rightarrow$  constant regions same as input
- Amount of smoothing proportional to mask size
- Remove “high-frequency” components; “low-pass” filter

- Derivatives

- Opposite signs used to get high response in regions of high contrast
- Sum to 0  $\rightarrow$  no response in constant regions
- High absolute value at points of high contrast

- Filters act as templates

- Highest response for regions that “look the most like the filter”
- Dot product as correlation

# Summary so far

- Image gradients
- Seam carving – gradients as “energy”
- Gradients → edges and contours
- Template matching
  - Image patch as a filter

# Next

- Edge detection and matching
  - process the image gradient to find curves/contours
  - comparing contours
- Binary image analysis
  - blobs and regions



# Motivation



Fig. 1. Examples of two handwritten digits. In terms of pixel-to-pixel comparisons, these two images are quite different, but to the human observer, the shapes appear to be similar.

# Chamfer distance

- Average distance to nearest feature


$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$

$I$  = Set of points in image

$T$  = Set of points on (shifted) template

$d_I(t)$  = Minimum distance between point  $t$   
and some point in  $I$

# Chamfer distance



$$D_{chamfer}(T, I) \equiv \frac{1}{|I|} \sum_{t \in I} d_I(t)$$


# Chamfer distance

- Average distance to nearest feature

$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$

*How is the measure different than just filtering with a mask having the shape points?*

*How expensive is a naïve implementation?*



Edge image

# Distance transform

*Image features (2D)*


*Distance Transform*

1	0	1	2	3	4	3	2
1	0	1	2	3	3	2	1
1	0	1	2	3	2	1	0
1	0	0	1	2	1	0	1
2	1	1	2	1	0	1	2
3	2	2	2	1	0	1	2
4	3	3	2	1	0	1	2
5	4	4	3	2	1	0	1


**Distance Transform** is a function  $D(\cdot)$  that for each image pixel  $p$  assigns a non-negative number  $D(p)$  corresponding to distance from  $p$  to the nearest feature in the image  $I$

Features could be edge points, foreground points,...


# Distance transform



original



edges



distance transform

Value at  $(x,y)$  tells how far that position is from the nearest edge point (or other binary image structure)

```
>> help bwdist
```

# Distance transform (1D)

Two pass  $O(n)$  algorithm for 1D  $L_1$  norm

1. Initialize: For all  $j$


$$D[j] \leftarrow 1_{\mathbf{P}}[j]$$

// 0 if  $j$  is in  $\mathbf{P}$ , infinity otherwise

# Distance Transform (2D)

- 2D case analogous to 1D
  - Initialization
  - Forward and backward pass
    - Fwd pass finds closest above and to left
    - Bwd pass finds closest below and to right

-	1
1	0
0	1
1	-



$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	$\infty$	$\infty$	$\infty$

$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	1	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	$\infty$	$\infty$	$\infty$

$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	1	2
$\infty$	0	1	2
$\infty$	1	2	3


2	1	2	3
1	0	1	2
1	0	1	2
2	1	2	3



# Chamfer distance

- Average distance to nearest feature

$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$




Edge image



Distance transform image

# Chamfer distance




Edge image

Distance transform image

# Chamfer distance: properties

- Sensitive to scale and rotation
- Tolerant of small shape changes, clutter
- Need large number of template shapes
- Inexpensive way to match shapes

# Chamfer matching system



- Gavrilă et al.  
[http://gavrila.net/Research/Chamfer\\_System/chamfer\\_system.html](http://gavrila.net/Research/Chamfer_System/chamfer_system.html)

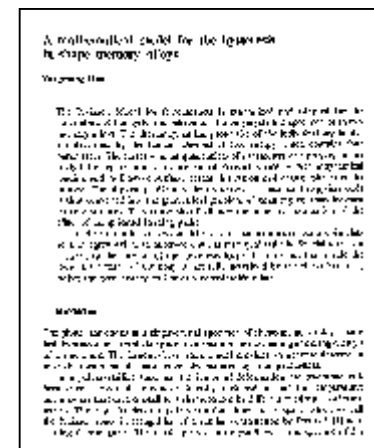


# Today

- Edge detection and matching
  - process the image gradient to find curves/contours
  - comparing contours
- Binary image analysis
  - blobs and regions

# Binary images



0	0	0	0	0
1	1	1	1	1
2	2	2	2	2
3	3	3	3	3
4	4	4	4	4
5	5	5	5	5
6	6	6	6	6
7	7	7	7	7
8	8	8	8	8
9	9	9	9	9




# Binary image analysis: basic steps

- Convert the image into binary form
  - Thresholding
- Clean up the thresholded image
  - Morphological operators
- Extract separate blobs
  - Connected components
- Describe the blobs with region properties

# Binary images

- Two pixel values
  - Foreground and background
  - Mark region(s) of interest

1	1	0	1	1	1	0	1
1	1	0	1	0	1	0	1
1	1	1	1	0	0	0	1
0	0	0	0	0	0	0	1
1	1	1	1	0	1	0	1
0	0	0	1	0	1	0	1
1	1	0	1	0	0	0	1
1	1	0	1	0	1	1	1





# Thresholding

- Given a grayscale image or an intermediate matrix  $\rightarrow$  threshold to create a binary output.

Example: edge detection



Gradient magnitude



```
fg_pix = find(gradient_mag > t);
```

Looking for pixels where gradient is strong.

# Thresholding

- Given a grayscale image or an intermediate matrix  $\rightarrow$  threshold to create a binary output.


Example: background subtraction



-



=




Looking for pixels that differ significantly from the “empty” background.

```
fg_pix = find(diff > t);
```

# Thresholding

- Given a grayscale image or an intermediate matrix → threshold to create a binary output.

Example: intensity-based detection



```
fg_pix = find(im < 65);
```

Looking for dark pixels

# Thresholding

- Given a grayscale image or an intermediate matrix  $\rightarrow$  threshold to create a binary output.

Example: color-based detection





```
fg_pix = find(hue > t1 & hue < t2);
```

Looking for pixels within a certain hue range.

# Issues

- What to do with “noisy” binary outputs?
  - Holes
  - Extra small fragments
- How to demarcate multiple regions of interest?
  - Count objects
  - Compute further features per object




# Morphological operators

- Change the shape of the foreground regions via intersection/union operations between a scanning structuring element and binary image.
- Useful to clean up result from thresholding
- Basic operators are:
  - Dilation
  - Erosion

# Dilation

- Expands connected components
- Grow features
- Fill holes




**Before dilation**




**After dilation**

# Erosion

- Erode connected components
- Shrink features
- Remove bridges, branches, noise



**Before erosion**




**After erosion**



# Structuring elements

- **Masks** of varying shapes and sizes used to perform morphology, for example:



- Scan mask across foreground pixels to transform the binary image

```
>> help strel
```

# Dilation vs. Erosion

At each position:

- **Dilation:** if current pixel is foreground, OR the structuring element with the input image.

# Example for Dilation (1D)


**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



1	1	1
---	---	---

**Structuring Element**



**Output Image**

1									
---	--	--	--	--	--	--	--	--	--

$$g(x) = f(x) \oplus SE$$

# Example for Dilation

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---




**Output Image**

1	1								
---	---	--	--	--	--	--	--	--	--

# Example for Dilation


**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**

1	1	0							
---	---	---	--	--	--	--	--	--	--

# Example for Dilation


**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---




**Output Image**

1	1	0	1						
---	---	---	---	--	--	--	--	--	--

# Example for Dilation

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**


1	1	0	1	1					
---	---	---	---	---	--	--	--	--	--



# Example for Dilation

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**

1	1	0	1	1	1				
---	---	---	---	---	---	--	--	--	--

# Example for Dilation

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---




**Output Image**

1	1	0	1	1	1	1			
---	---	---	---	---	---	---	--	--	--

# Example for Dilation


**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---




**Output Image**

1	1	0	1	1	1	1	1		
---	---	---	---	---	---	---	---	--	--

# Example for Dilation


**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**

1	1	0	1	1	1	1	1	1	
---	---	---	---	---	---	---	---	---	--

Note that the object gets bigger and holes are filled.

```
>> help imdilate
```

# 2D example for dilation

1	1	1	1	1	1	1	
			1	1	1	1	
			1	1	1	1	
		1	1	1	1	1	
			1	1	1	1	
		1	1				

(a) Binary image **B**

1	1	1
1	1	1
1	1	1

(b) Structuring element **S**

1	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1
	1	1	1	1	1	1	1
	1	1	1	1	1	1	1
	1	1	1	1	1	1	1
	1	1	1	1	1	1	1
	1	1	1	1			

(c) Dilation  $\mathbf{B} \oplus \mathbf{S}$

# Dilation vs. Erosion

At each position:

- **Dilation:** if **current pixel** is foreground, OR the structuring element with the input image.
- **Erosion:** if **every pixel** under the structuring element's nonzero entries is foreground, OR the current pixel with S.

# Example for Erosion (1D)

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



1	1	1
---	---	---

**Structuring Element**



**Output Image**

0									
---	--	--	--	--	--	--	--	--	--

$$g(x) = f(x) \ominus SE$$



# Example for Erosion (1D)

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**


0	0								
---	---	--	--	--	--	--	--	--	--

$$g(x) = f(x) \ominus SE$$

# Example for Erosion


**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**

0	0	0							
---	---	---	--	--	--	--	--	--	--

# Example for Erosion


**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---




**Output Image**

0	0	0	0						
---	---	---	---	--	--	--	--	--	--

# Example for Erosion

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**

0	0	0	0	0					
---	---	---	---	---	--	--	--	--	--

# Example for Erosion

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---




**Output Image**

0	0	0	0	0	1				
---	---	---	---	---	---	--	--	--	--

# Example for Erosion

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**

0	0	0	0	0	1	0			
---	---	---	---	---	---	---	--	--	--

# Example for Erosion

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**

0	0	0	0	0	1	0	0		
---	---	---	---	---	---	---	---	--	--



# Example for Erosion

**Input image**

1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---



**Structuring Element**

1	1	1
---	---	---



**Output Image**


0	0	0	0	0	1	0	0	0	
---	---	---	---	---	---	---	---	---	--

# Example for Erosion

**Input image**

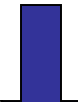
1	0	0	0	1	1	1	0	1	1
---	---	---	---	---	---	---	---	---	---

**Structuring Element**



1	1	1
---	---	---

**Output Image**



0	0	0	0	0	1	0	0	0	1
---	---	---	---	---	---	---	---	---	---

Note that the object gets smaller

```
>> help imerode
```

# 2D example for erosion

1	1	1	1	1	1	1	
			1	1	1	1	
			1	1	1	1	
		1	1	1	1	1	
			1	1	1	1	
		1	1				

(a) Binary image  $\mathbf{B}$

1	1	1
1	1	1
1	1	1


(b) Structuring element  $\mathbf{S}$

				1	1		
				1	1		
				1	1		


(d) Erosion  $\mathbf{B} \ominus \mathbf{S}$

# Opening

- Erode, then dilate
- Remove small objects, keep original shape




Before opening




After opening

# Closing

- Dilate, then erode
- Fill holes, but keep original shape




Before closing



After closing

# Issues

- What to do with “noisy” binary outputs?
  - Holes
  - Extra small fragments
- How to demarcate multiple regions of interest?
  - Count objects
  - Compute further features per object



# Connected components



- Identify distinct regions of “connected pixels”

1	1	0	1	1	1	0	1
1	1	0	1	0	1	0	1
1	1	1	1	0	0	0	1
0	0	0	0	0	0	0	1
1	1	1	1	0	1	0	1
0	0	0	1	0	1	0	1
1	1	0	1	0	0	0	1
1	1	0	1	0	1	1	1

a) binary image

1	1	0	1	1	1	0	2
1	1	0	1	0	1	0	2
1	1	1	1	0	0	0	2
0	0	0	0	0	0	0	2
3	3	3	3	0	4	0	2
0	0	0	3	0	4	0	2
5	5	0	3	0	0	0	2
5	5	0	3	0	2	2	2


b) connected components labeling




c) binary image and labeling, expanded for viewing

# Connectedness

- Defining which pixels are considered neighbors



**4-connected**




**8-connected**



# Connected components

- We'll consider a sequential algorithm that requires only 2 passes over the image.
- **Input:** binary image
- **Output:** "label" image, where pixels are numbered per their component
- Note: foreground here is denoted with black pixels.



# Sequential connected components

- Labeling a pixel only requires to consider its prior and superior neighbors.
- It depends on the type of connectivity used for foreground (4-connectivity here).

Same object



(a)



(b)



(c)

New object



(d)


What happens in these cases?






# Sequential connected components

- Labeling a pixel only requires to consider its prior and superior neighbors.
- It depends on the type of connectivity used for foreground (4-connectivity here).


Same object



New object




What happens in these cases?



# Sequential connected components

- Labeling a pixel only requires to consider its prior and superior neighbors.
- It depends on the type of connectivity used for foreground (4-connectivity here).





Same object



New object



What happens in these cases?



# Sequential connected components

- Process the image from left to right, top to bottom.

1. If the next pixel to process is 1-pixel:

**Already processed**

1. If only one of its neighbors (superior or left) is 1-pixel, copy its label.

2. If both are, and have the same label, copy it.

3. If they have different labels:


**superior? smallest?**

1. Copy the label from the prior.

2. Reflect the change in the table of equivalences.

4. Otw, assign a new label.



2. More pixels? Go to step 1.




$\{1, 3, 4, 5, \}$      $2, 7\}$   
 $\{6, 8\}$

- Re-label with the smallest of equivalent labels.
- Pixels of the same segment always have the same label.

# Connected components





connected  
components  
of 1's from  
thresholded  
image



connected  
components  
of cluster  
labels

# Region properties

- Given connected components, can compute simple features per blob, such as:
  - Area (num pixels in the region)
  - Centroid (average x and y position of pixels in the region)
  - Bounding box (min and max coordinates)
  - Circularity (ratio of mean dist. to centroid over std)



# Binary image analysis: basic steps (recap)

- Convert the image into binary form
  - Thresholding
- Clean up the thresholded image
  - Morphological operators
- Extract separate blobs
  - Connected components
- Describe the blobs with region properties



# Matlab

- `N = hist(Y,M)`
- `L = bwlabel (BW,N) ;`
- `STATS = regionprops (L,PROPERTIES) ;`
  - `'Area'`
  - `'Centroid'`
  - `'BoundingBox'`
  - `'Orientation', ...`
- `IM2 = imerode (IM,SE) ;`
- `IM2 = imdilate (IM,SE) ;`
- `IM2 = imclose (IM, SE) ;`
- `IM2 = imopen (IM, SE) ;`

# Example using binary image analysis: OCR



The screenshot shows the reCAPTCHA website interface. On the left is a navigation menu with links: HOME, WHAT IS reCAPTCHA, DIGITIZATION ACCURACY, WHAT IS A CAPTCHA, SECURITY, GET reCAPTCHA, MY ACCOUNT, EMAIL PROTECTION, and RESOURCES. The main heading is "Digitizing Books One Word at a Time". The central task area shows two distorted words, "vibratos" and "ing", with a text input field and a "Submit" button. To the right of the input field are icons for refreshing, audio, and help, along with the reCAPTCHA logo and the text "stop spam. read books.". Below the task area, a message states: "The words above come from scanned books. By typing them, you help to digitize old texts." At the bottom, a paragraph explains that reCAPTCHA is a free service for digitizing books and old-time radio shows, and provides a link to "our paper" in Science. Another paragraph defines CAPTCHA as a program to distinguish humans from computers, noting its use in preventing spam.

reCAPTCHA™

Digitizing Books One Word at a Time

→ HOME  
→ WHAT IS reCAPTCHA  
DIGITIZATION ACCURACY  
WHAT IS A CAPTCHA  
SECURITY  
→ GET reCAPTCHA  
→ MY ACCOUNT  
→ EMAIL PROTECTION  
→ RESOURCES

Type the two words:

Submit

The words above come from scanned books.  
By typing them, you help to digitize old texts.


reCAPTCHA™ stop spam. read books.

reCAPTCHA is a free CAPTCHA service that helps to digitize books, newspapers and old time radio shows. Check out [our paper](#) in Science about it (or read more below).

A [CAPTCHA](#) is a program that can tell whether its user is a human or a computer. You've probably seen them — colorful images with distorted text at the bottom of Web registration forms. CAPTCHAs are used by many websites to prevent abuse from "bots," or automated programs usually written to generate spam. No computer program can read distorted text as well as humans can, so bots cannot navigate sites protected by CAPTCHAs.

[Luis von Ahn et al. <http://recaptcha.net/learnmore.html>]


# Example using binary image analysis: segmentation of a liver




Threshold



Extract Largest  
Region



Region Filling



Extract Largest  
Region



# Binary images

- Pros
  - Can be fast to compute, easy to store
  - Simple processing techniques available
  - Lead to some useful compact shape descriptors
- Cons
  - Hard to get “clean” silhouettes
  - Noise common in realistic scenarios
  - Can be too coarse of a representation
  - Not 3d

# New concepts today

- Gradients -> edge maps
- Template matching
- Chamfer distance
- Distance transform
- Erosion/dilation for binary images
- Connected components
- Region properties

# Assignments

- 了解概念：二值图像连通区域。编写“种子填充法” 标记二值图像的连通区域的伪代码，并讨论其算法复杂度。
- 编程实现二值图像连通区域标记，并计算其面积等统计量。